

ctf

October 15, 2019

```
In [7]: import json
import pandas as pd
import seaborn as sns; sns.set()
import babeltrace as bl
```

```
In [8]: tr = bl.TraceCollection()
tr.add_trace('sample_traces/jsoo/ctf', "ctf")
```

```
events = []
counters = []
allocs = []
```

```
ENTRY = 0
EXIT = 1
COUNTER = 2
ALLOC = 3
```

*# babeltrace is an evil binding, do not put Event objects directly in a list, need to use
Events objects are actually references to the head of the iterator, so trying to use
the scope of the iterator will results in having the same objects, that is, the last
`tr.events()`*

```
for event in tr.events:
    id = event['id']
    ev = {}
    ev['timestamp'] = event['timestamp']
    if id == ENTRY or id == EXIT:
        ev['phase'] = event['phase']
        ev['id'] = id
        events.append(ev)
    if id == COUNTER:
        ev['name'] = event['kind']
        ev['count'] = event['count']
        counters.append(ev)
    if id == ALLOC:
        ev['bucket'] = event['bucket']
        ev['count'] = event['count']
```

```

        allocs.append(ev)

    timeline_set = {}
    timeline = []

    for event in events:
        phase = event['phase']
        if event['id'] == ENTRY:
            if phase in timeline_set:
                print("overlapping entry for phase: %s, timestamp %d" % (phase, event['timestamp']))
            else:
                timeline_set[phase] = event
        elif event['id'] == EXIT:
            if phase in timeline_set:
                ev_entry = timeline_set[phase]
                item = {
                    'ts_begin': ev_entry['timestamp'],
                    'ts_end': event['timestamp'],
                    'dur': event['timestamp'] - ev_entry['timestamp'],
                    'phase': phase
                }
                timeline.append(item)
                del timeline_set[phase]
            else:
                print("got an exit event with no entry: %s, timestamp %d" % (phase, event['timestamp']))

    for key in timeline_set:
        print("got an entry event with no exit: %s, timestamp: %d" % (key, timeline_set[key]['timestamp']))

    # FIXME: none of these should be tolerable besides impromptu program shutdown (segfault)

    got an exit event with no entry: major/mark_final, timestamp 4825128
    overlapping entry for phase: major/mark_roots, timestamp 11983480
    got an exit event with no entry: major/mark_final, timestamp 13656935
    overlapping entry for phase: major/mark_roots, timestamp 21288580
    got an exit event with no entry: major/mark_final, timestamp 26077514
    overlapping entry for phase: major/mark_roots, timestamp 34357892
    got an exit event with no entry: major/mark_final, timestamp 40564907
    overlapping entry for phase: major/mark_roots, timestamp 49507346
    got an exit event with no entry: major/mark_main, timestamp 59133890
    got an exit event with no entry: major/mark_final, timestamp 61191983
    overlapping entry for phase: major/mark_roots, timestamp 76844490
    got an exit event with no entry: major/mark_final, timestamp 95465765
    overlapping entry for phase: major/mark_roots, timestamp 118788736
    got an exit event with no entry: major/mark_final, timestamp 125912431
    overlapping entry for phase: major/mark_roots, timestamp 194698284
    got an exit event with no entry: major/mark_final, timestamp 211612669
    overlapping entry for phase: major/mark_roots, timestamp 250176029

```

got an exit event with no entry: major/mark_final, timestamp 283058871
overlapping entry for phase: major/mark_roots, timestamp 322517850
got an exit event with no entry: major/mark_final, timestamp 377884343
overlapping entry for phase: major/mark_roots, timestamp 427734240
got an exit event with no entry: major/mark_final, timestamp 507464932
overlapping entry for phase: major/mark_roots, timestamp 580211312
got an exit event with no entry: major/mark_final, timestamp 690870591
overlapping entry for phase: major/mark_roots, timestamp 792030298
got an exit event with no entry: major/mark_final, timestamp 918199661
overlapping entry for phase: major/mark_roots, timestamp 1062123183
got an exit event with no entry: major/mark_final, timestamp 1170517771
overlapping entry for phase: major/mark_roots, timestamp 1300095797
got an exit event with no entry: major/mark_final, timestamp 1493961343
overlapping entry for phase: major/mark_roots, timestamp 1635330632
got an exit event with no entry: major/mark_final, timestamp 2099296663
overlapping entry for phase: major/mark_roots, timestamp 2490591161
got an exit event with no entry: major/mark_final, timestamp 2728797746
overlapping entry for phase: major/mark_roots, timestamp 2978264669
got an exit event with no entry: major/mark_final, timestamp 3365051615
overlapping entry for phase: major/mark_roots, timestamp 3752049014
got an exit event with no entry: major/mark_final, timestamp 4213815712
overlapping entry for phase: major/mark_roots, timestamp 4773473417
got an exit event with no entry: major/mark_final, timestamp 4867381081
overlapping entry for phase: major/mark_roots, timestamp 5244432043
got an exit event with no entry: major/mark_final, timestamp 5384739792
overlapping entry for phase: major/mark_roots, timestamp 5591135483
got an exit event with no entry: major/mark_final, timestamp 5738045294
overlapping entry for phase: major/mark_roots, timestamp 6097403181
got an exit event with no entry: major/mark_final, timestamp 6234878457
overlapping entry for phase: major/mark_roots, timestamp 6629145536
got an exit event with no entry: major/mark_final, timestamp 6833722761
overlapping entry for phase: major/mark_roots, timestamp 7631653769
got an exit event with no entry: major/mark_final, timestamp 7659529071
overlapping entry for phase: major/mark_roots, timestamp 7721659376
got an exit event with no entry: major/mark_final, timestamp 7821715956
overlapping entry for phase: major/mark_roots, timestamp 7885020297
got an exit event with no entry: major/mark_final, timestamp 8029802278
overlapping entry for phase: major/mark_roots, timestamp 8115721144
got an exit event with no entry: major/mark_final, timestamp 8231446913
overlapping entry for phase: major/mark_roots, timestamp 8422551363
got an exit event with no entry: major/mark_final, timestamp 8594427040
overlapping entry for phase: major/mark_roots, timestamp 8903335494
got an exit event with no entry: major/mark_final, timestamp 8953525683
overlapping entry for phase: major/mark_roots, timestamp 9261368236
got an exit event with no entry: major/mark_final, timestamp 9364798938
overlapping entry for phase: major/mark_roots, timestamp 9473523337
got an exit event with no entry: major/mark_final, timestamp 9587774617
overlapping entry for phase: major/mark_roots, timestamp 9707670096

```

got an exit event with no entry: major/mark_final, timestamp 9788359022
overlapping entry for phase: major/mark_roots, timestamp 9960560125
got an exit event with no entry: major/mark_final, timestamp 10073951034
overlapping entry for phase: major/mark_roots, timestamp 10174445531
got an exit event with no entry: major/mark_final, timestamp 10289754054
overlapping entry for phase: major/mark_roots, timestamp 10574287546
got an exit event with no entry: major/mark_final, timestamp 10656694697
overlapping entry for phase: major/mark_roots, timestamp 11205460968
got an exit event with no entry: major/mark_final, timestamp 11351559488
overlapping entry for phase: major/mark_roots, timestamp 11635824955
got an exit event with no entry: major/mark_final, timestamp 11837322064
overlapping entry for phase: major/mark_roots, timestamp 12026483849
got an exit event with no entry: major/mark_final, timestamp 12192500510
overlapping entry for phase: major/mark_roots, timestamp 12381684475
got an exit event with no entry: major/mark_final, timestamp 12487036085
overlapping entry for phase: major/mark_roots, timestamp 12665799016
got an exit event with no entry: major/mark_final, timestamp 12841273654
overlapping entry for phase: major/mark_roots, timestamp 12971735084
got an exit event with no entry: major/mark_final, timestamp 13105502885
overlapping entry for phase: major/mark_roots, timestamp 13309786748
got an exit event with no entry: major/mark_final, timestamp 13442618466
overlapping entry for phase: major/mark_roots, timestamp 13566654657
got an exit event with no entry: major/mark_final, timestamp 13714670984
overlapping entry for phase: major/mark_roots, timestamp 13869053057
got an exit event with no entry: major/mark_final, timestamp 14005690719
overlapping entry for phase: major/mark_roots, timestamp 14362559551
got an exit event with no entry: major/mark_final, timestamp 14567846636
overlapping entry for phase: major/mark_roots, timestamp 14713706140
got an exit event with no entry: major/mark_final, timestamp 14952685362
overlapping entry for phase: major/mark_roots, timestamp 15266114344
got an exit event with no entry: major/mark_final, timestamp 15446310460
overlapping entry for phase: major/mark_roots, timestamp 15735715071
got an exit event with no entry: major/mark_final, timestamp 15964232283
overlapping entry for phase: major/mark_roots, timestamp 18917065824
got an exit event with no entry: major/mark_final, timestamp 19156692461
overlapping entry for phase: major/mark_roots, timestamp 19333929041
got an exit event with no entry: major/mark_final, timestamp 19547936648
overlapping entry for phase: major/mark_roots, timestamp 20084463630
got an entry event with no exit: major/mark_main, timestamp: 19546706309

```

In [9]: *# From here, refer to <https://github.com/Engil/gc-tracing/blob/master/metadata> to know
the various phases and counters available.*

```

def counters_by_name(name): return [ev for ev in counters if ev['name'] == name]
def events_by_name(name): return [ev for ev in timeline if ev['phase'] == name]


s = pd.DataFrame(counters_by_name('minor/promoted'))

```

```
s['timestamp'] = s['timestamp'].astype('timedelta64[ns]')
```

```
sns.lineplot(x="timestamp", y="count", data=s)
```

Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd9676b0e48>




In [10]: *# Visiting allocations during the lifetime of the program, by buckets*


```
s = pd.DataFrame(allocs)
```

```
ax = sns.catplot(x='timestamp', y='count', hue='bucket', data=s)  
ax.set(xticks=[])
```

Out[10]: <seaborn.axisgrid.FacetGrid at 0x7fd973daaf28>



```
In [11]: # minor and major dispatches, duration on the Y axis
s = pd.DataFrame([ev for ev in timeline if ev['phase'] == 'major' or ev['phase'] == 'minor'])
ax = sns.scatterplot(x='ts_begin', y='dur', hue='phase', data=s)
```



```
In [12]: def gc_phase_quantile(name):
          s = pd.DataFrame(events_by_name(name))
          #s['timestamp'] = s['timestamp'].astype('timedelta64[ns]')
          return (s['dur'].quantile([.1, .25, .50, .75, .95, .99]))
```

```
In [13]: gc_phase_quantile('major')
```

```
Out[13]: 0.10      8234.40
         0.25     162240.00
         0.50     477792.00
         0.75     1172755.50
         0.95     3855911.30
         0.99     10411310.18
         Name: dur, dtype: float64
```

```
In [14]: gc_phase_quantile('minor')
```

```
Out[14]: 0.10      77132.40
         0.25     149607.00
         0.50     396085.00
         0.75     692175.00
         0.95     2578264.40
         0.99     3257261.36
         Name: dur, dtype: float64
```

```
In [15]: gc_phase_quantile('major/sweep')
```

```
Out[15]: 0.10      2395.20  
         0.25     119395.50  
         0.50     350880.50  
         0.75     627666.75  
         0.95     2317725.45  
         0.99     4270188.74  
         Name: dur, dtype: float64
```